

SYSTEM DESIGN

"CONTINUE WATCHING"

in OTT platforms

{ Discussion focused on Master Architecture }

What is "Continue watching" - Functional Req.

- ① You don't have to remember wherever you last left off the video.
- ② Experience across the devices - Sync the content watched across the devices user is logged in.

Due to #2 above, we can't use In-app Cache, the data should be put somewhere at central location to sync across the devices.

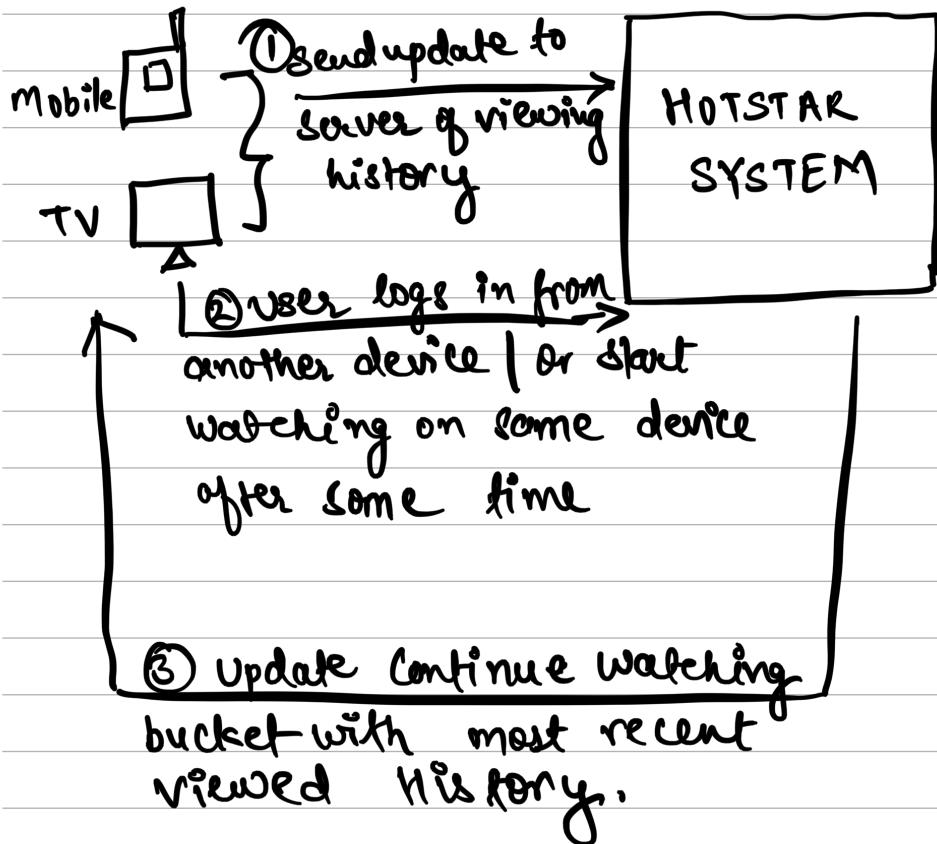
Non Functional Requirements :

- ① Low latency - should be synced across the devices in no time.
- ② Availability
- ③ Durability
- ④ Eventual consistency
↓
related to #1
- ⑤ less cost - Operations & Money.

User Statistics (as per Internet)

50M subscribers → users watch avg 1B minutes of video/day
300M active users → This requires ~200GB/day storage for continue watching feature

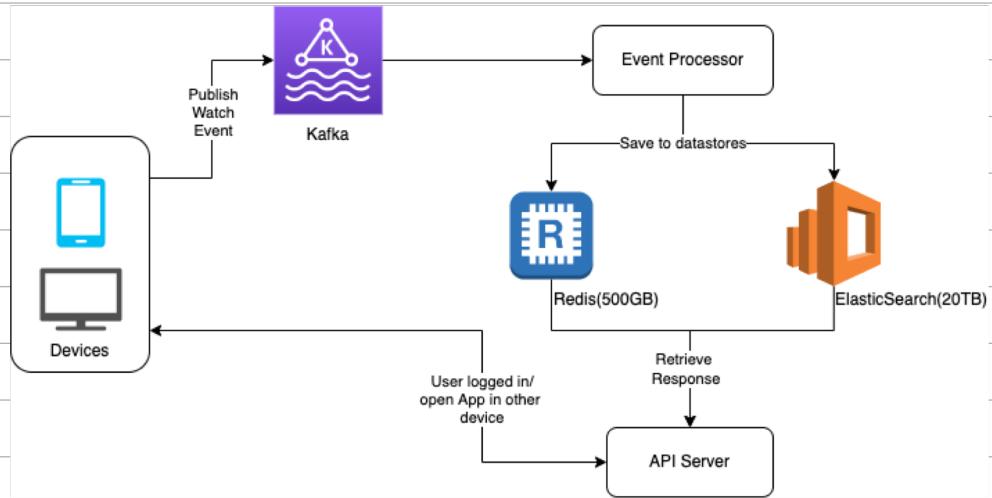
Very High Level View of Architecture →



- Hotstar system should consist of database(s) which scales for 300M plus active users and maintains the watch history of every user.
- The database will definitely be write heavy.
- The read latency should be minimum - user's "Continue watching" bucket should be updated as soon they login (open the application).
- Database should scale in heavy traffic times, request volume increases by 10 to 20 times within a minute.

Architecture 1

Use combination of Redis and Elasticsearch for data storage and retrieval.



Database Model →

Redis

userId	hist<content Id>
mandeep	{movie1, movie2}

Elastic Search

userId-content-id: {

```

"field1": xyz
"field2": pqr
  
```

}

→ Redis maintains list of content for a user.

→ Elastic Search Index contains metadata for user-contentId.
Example →

till what timestamp video is watched
which episode of series is in progress.

Bottlenecks in this design

① Redis latency was upto the mark, but there is requirement to horizontally scale the cluster as data size grew.
↳ increased cost every 3-4 months.

② ES latency ~200 milliseconds

↳ avg cost of Elastic Search was at higher end.
↳ its just a user experience feature

↳ Node maintenance issues and manual effort requirement to resolve issues.

↳ We can minimize this keeping cluster size small - Watch "How Discord Index Billions of messages" for complete analysis.

③ Two data stores for same use-case. - becomes difficult to maintain two data stores with different code base and query pattern at high scale.

↳ Extra administration cost.

Architecture Improvements →

- ① think if we can solve this with single database.
- ② Notstar choose ScyllaDB as a replacement of Redis and ES cluster

→ how latency for read/write
→ Operational simplicity
→ how cost - compared to other NOSQL Solutions available in Cloud.

if it is based on Cassandra

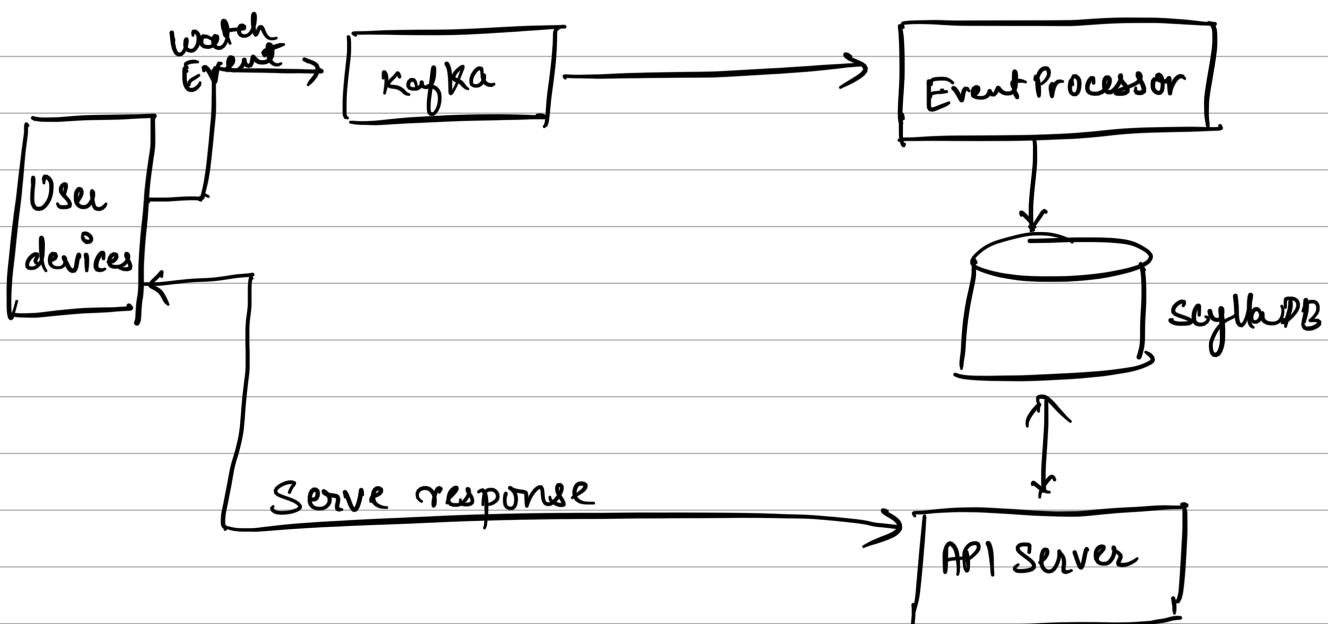
Scylla DB Model

- ① User Table

userId (PK)	List <content-id>
-------------	-------------------

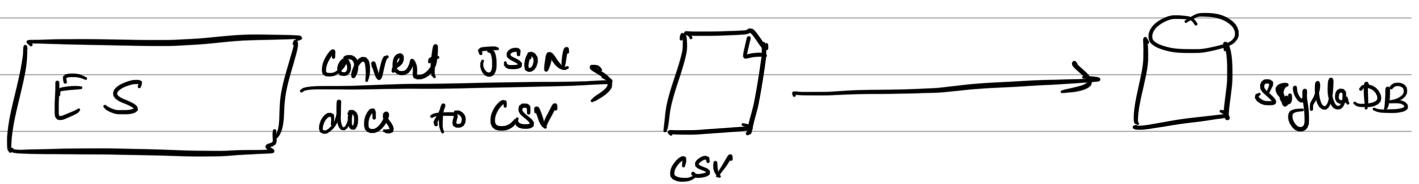
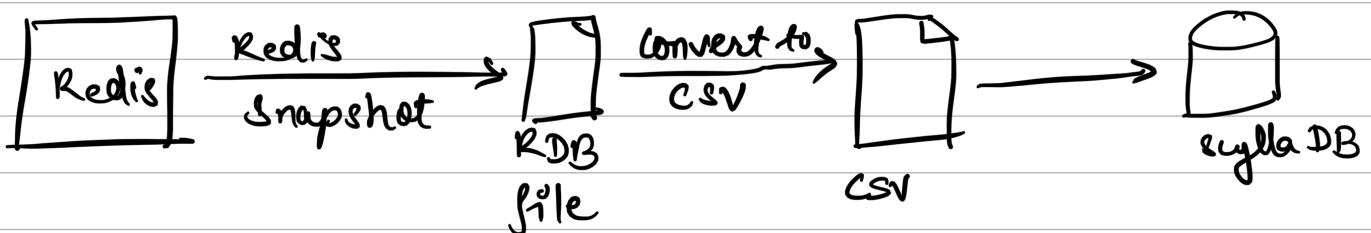
- ② User Content Table

userId (PK)	ContentId (SK)	Timestamp	field 3	field 4
↑ Primary Key	↑ secondary key (clustering key)			

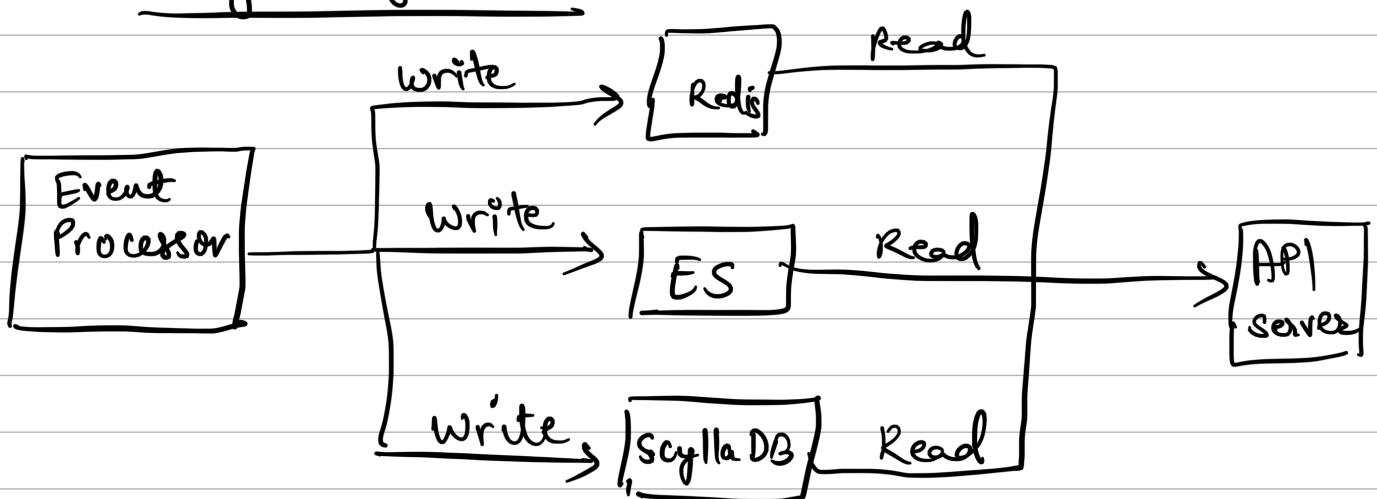


Migration from old to new with
 { NO DOWNTIME }

- ① Move old data from Redis|ES to ScyllaDB
- ② Run both system in parallel for confidence building
- ③ Remove old system.



During Migration :



Subscribe for more such content.

YouTube Channel - Ms Deep Singh

Happy Learning 😊