

How Discord Indexes Billions of msgs?

Subscribe YouTube channel - 'MsDeep Singh' for more such content.

Requirements

- ① Cost Effective - Search is accessory feature so price should not exceed actual storage.
- ② Fast & Intuitive
- ③ Self healing
- ④ Linearly scalable
- ⑤ Lazily Indexed - don't index msgs unless someone attempts to search them atleast once.

Finalized Solution - Elastic Search

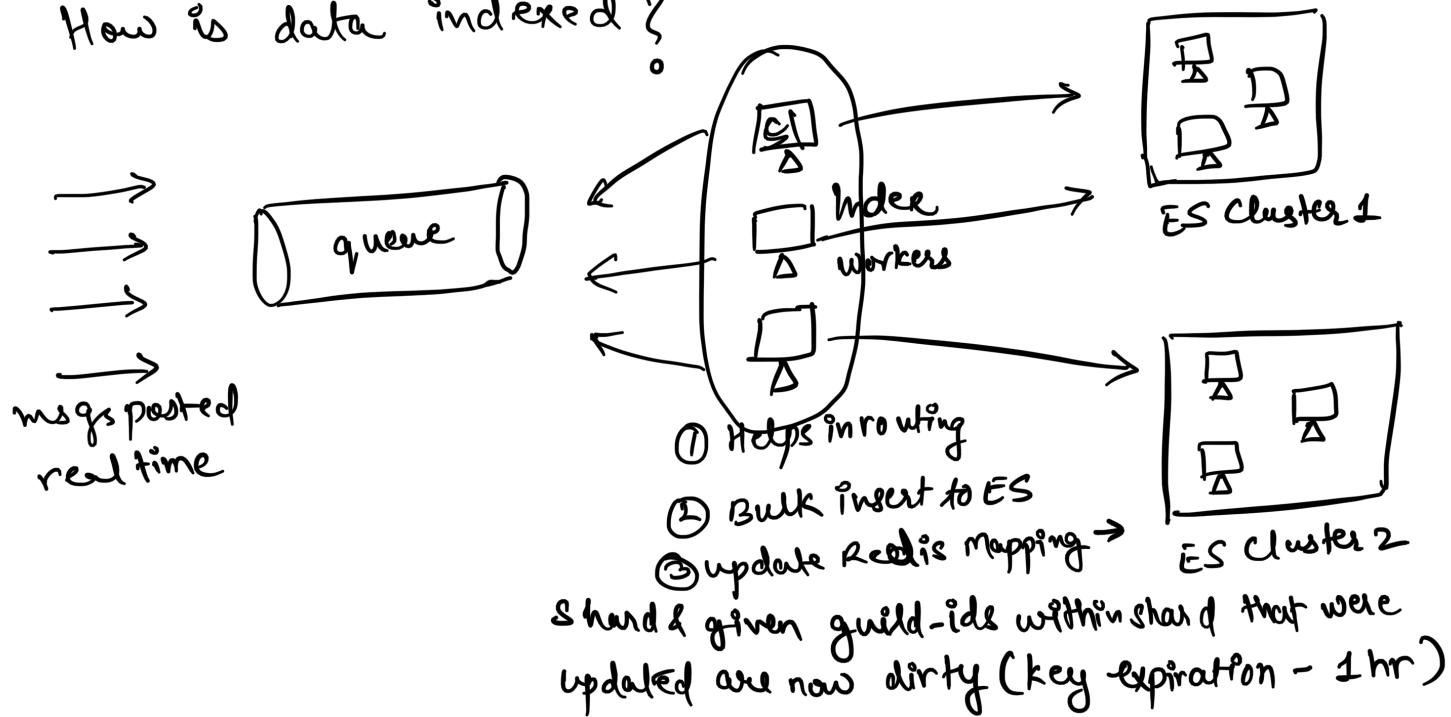
Why?

- ① No extra setup required for node discovery. Solr requires Zookeep
- ② ES supports automatic shard rebalancing on addition/removal of nodes.
- ③ Structured query DSL built in.
- ④ Previous work experience with ES.

How can you avoid large ES clusters as msg data grows?

- ① Delegate sharding and routing logic to application layer - this will help to index msgs into pool of smaller ES clusters.
- ② Multiple clusters also help to reduce blast radius.

How is data indexed?



Mapping of Discord Server's message to ES cluster + Index

It comes in two layers -

- ① Persistent Shard mapping - put in Cassandra.
- ② Shard mapping cache - while ingesting msgs on workers, querying Cassandra for **Shard** is slow operation. So cache these mappings in Redis.

{ES cluster + index}

Don't confuse with ES shard.

Select shard to Discard sever's messages (when indexed first time)

Shard Allocator : ① Use sorted set in Redis, keep set of shards with score that represent their load.

- ② The shard with lowest load should be allocated next.
- ③ Score gets incremented on every new allocation.

Discovery of ES clusters and hosts within them -

etcd is used. — nodes in ES cluster can announce themselves onto etcd.

Indexing and mapping the Data

- ① ES cluster → consists of indices → contains no. of shards within it
- ↑
Shard is
lucene index.

- ② ES holds responsibility of data distribution within index to a shard belonging to that index.

↳ Discard has its own sharding logic in application layer (as described above in document)

- Ⓐ index contains one shard (no sharding required)
- Ⓑ index replication to one node.
- Ⓒ refresh interval - 60 minutes
- Ⓓ index contains single document type - message.

→ msg is transformed into bunch of fields containing metadata → index & search on this.

index template snippet → refer on discord blog

- message data is not duplicated in ES cluster, actual msg is fetched from cassandra.
 - ↳ saves additional disk space.
 - ↳ to highlight text matches, tokenizers and language analyzers are built into client application.
- No separate search service. Created a library as a wrapper of routing and querying logic to ES cluster.

why refresh interval is 60 minutes and not default 1sec

- ① 1sec refresh interval provides near time search ability.
- ② So every second for across the indices, ES flushes the in-memory buffer to Lucene Segment and opening the segment to make it searchable.

↳ ↑ get will require high CPU.

↳ More disk usage

↳ At less load time, ES merges massive amount of tiny segments into larger segments [which are space efficient]

Search Life cycle

- ① look up shard that needs to queried for guild-id
- ② Check Redis - if shard & guild-id is dirty.
- ③ If dirty, do a refresh of shard's ES index and mark entire shard as clean.
- ④ Execute search query & return results.

This helps in identifying if index should be refreshed. → since it is set to 1hr but if refresh is required prior to that for specific discord's server, it can be achieved.

What is guild-id - Helps in sloping a search to given server.

Guild - represents collection of users and channels.
From client perspective, guild means server.

How does search query works like?

SearchQuery
(guild-id = 12345678989 ,
content = "Hey Jake" ,
channel-ids = [1667089 , 38225979]
)

Happy Learning 😊

Subscribe for more

YouTube - MsDeep Singh
